

AI-assisted reporting with human approval.

A weekly Technology Operations reporting workflow turns scattered operational inputs into a leadership-ready draft, while I verify the evidence, priorities, and final voice before anything is shared.

SOURCES

DRAFT

REVIEW

APPROVED REPORT

How the reporting workflow operates

01

Collect the operating picture

Team recap submissions, meeting notes, Work Radar items, a prior finalized report, and live Jira queries are gathered as the source set.

02

Ask AI to shape the draft

The assistant organizes the source material into reporting sections, carries forward open priorities, and highlights gaps instead of inventing answers.

03

Verify before it leaves the room

I check source completeness, Jira counts, status and priority changes, tone, omissions, and whether each claim is safe to share.

04

Approve the leadership version

Only the reviewed and finalized report becomes the record. AI accelerates synthesis; accountability stays human.

AI does

- Structures a first draft from known inputs
- Carries forward priorities from prior reports
- Surfaces missing or ambiguous source areas
- Suggests concise leadership-reporting language

Billy decides

- Confirms every source was included
- Checks Jira counts and status changes
- Removes or rewrites anything sensitive
- Approves the final report and owns the decision

Bounded synthesis, verified output, human authority

01

No sensitive identifiers

The public case study avoids confidential project names, ticket keys, team names, screenshots, source data, credentials, internal systems, and private memory contents.

02

Evidence stays inspectable

The workflow treats source completeness and query results as checks, not assumptions.

03

AI cannot finalize

Drafting and synthesis are delegated. Approval, business judgment, and communication responsibility are not.

Latest experiment: Lumi / Hermes

Lumi/Hermes is a persistent AI partner runtime I use for real work. It can load context, invoke tools, edit files, run scheduled tasks, and verify outputs. The experiment is not replacing judgment; it is making the path from intent to tested artifact shorter.

1. Load relevant context
2. Use tools on real artifacts
3. Run verification
4. Return proof for human decision